

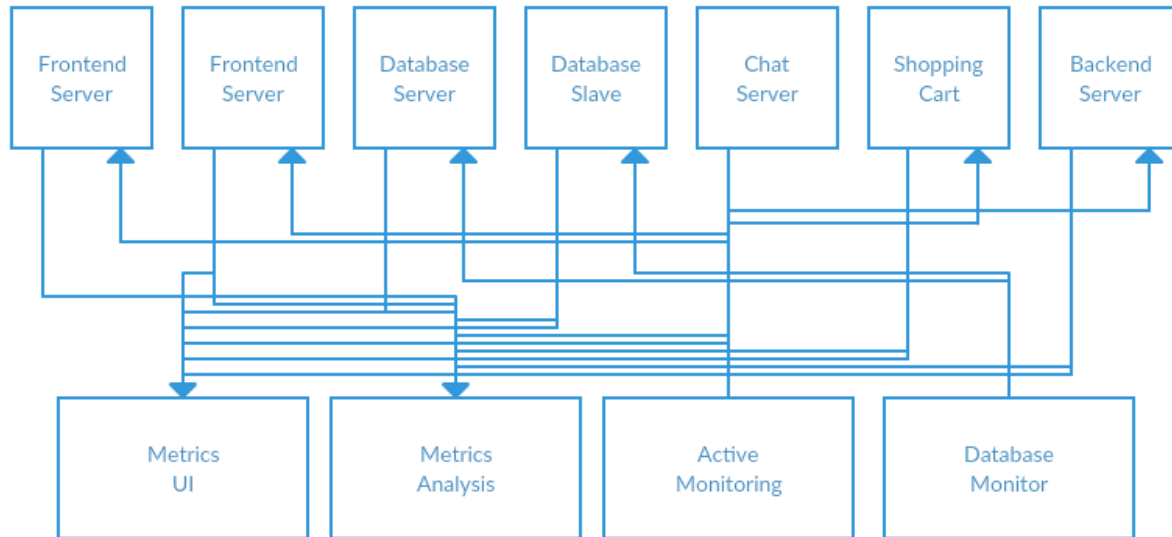


Kafka Introduction

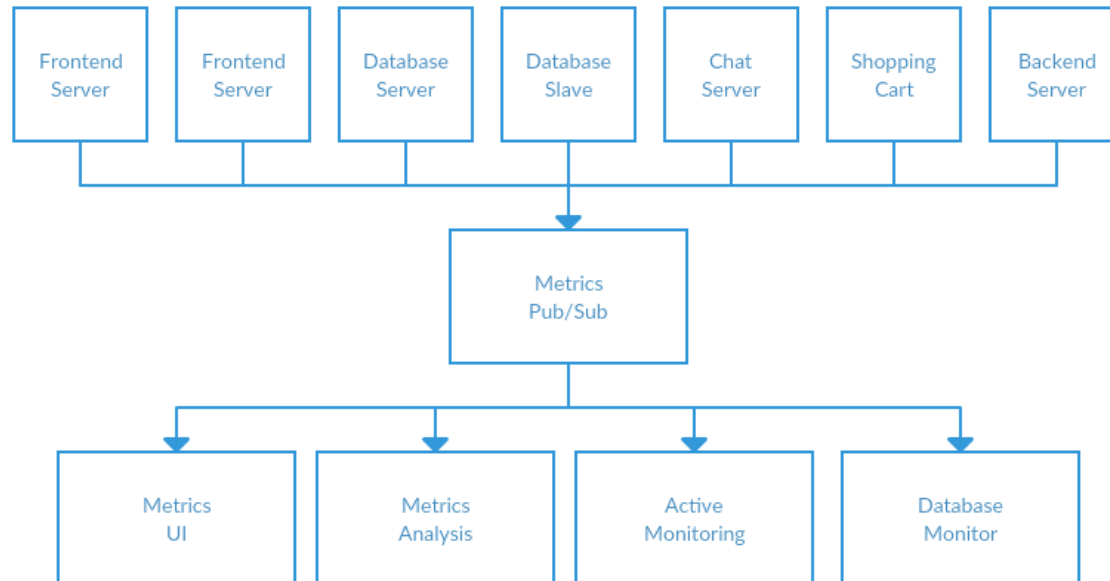
- ❑ Kafka is a leading general purpose ***publish-subscribe distributed messaging system***, which offers strong durability, scalability and fault-tolerance support.
- ❑ Designed for handling of ***real time activity*** stream data such as logs, metrics collections.
- ❑ Written in ***Scala***
- ❑ An apache project initially developed at ***LinkedIn*** at the year 2007 then to Apache in 2012 by Jay Kreps in the name of Franz Kafka
- ❑ It is **not specifically designed for Hadoop** rather Hadoop ecosystem is just be one of its possible consumers.

Why we need Kafka

Direct connections



Publish/Subscribe system





Rain

Producers



Dam

**Message
Queue**



River



Canal



Canal

Consumers



Why we need Kafka

- Kafka is Highly scalable very *easy to add large number of producers and consumers*
- Kafka handle spike
- Kafka also supports different consumption model
- Message durability is high in Kafka – *Persists*
- Integration support for Kafka with other frameworks are high

Applications



Clickstream



Geolocation



Web Data



Internet of Things




Docs, emails



Server logs

Ka&a High levelArchitecture

Kafka Architecture



Realtime Data Pipeline

OPERATIONAL
DATA HUB





STREAM + SQL
INTEGRATION

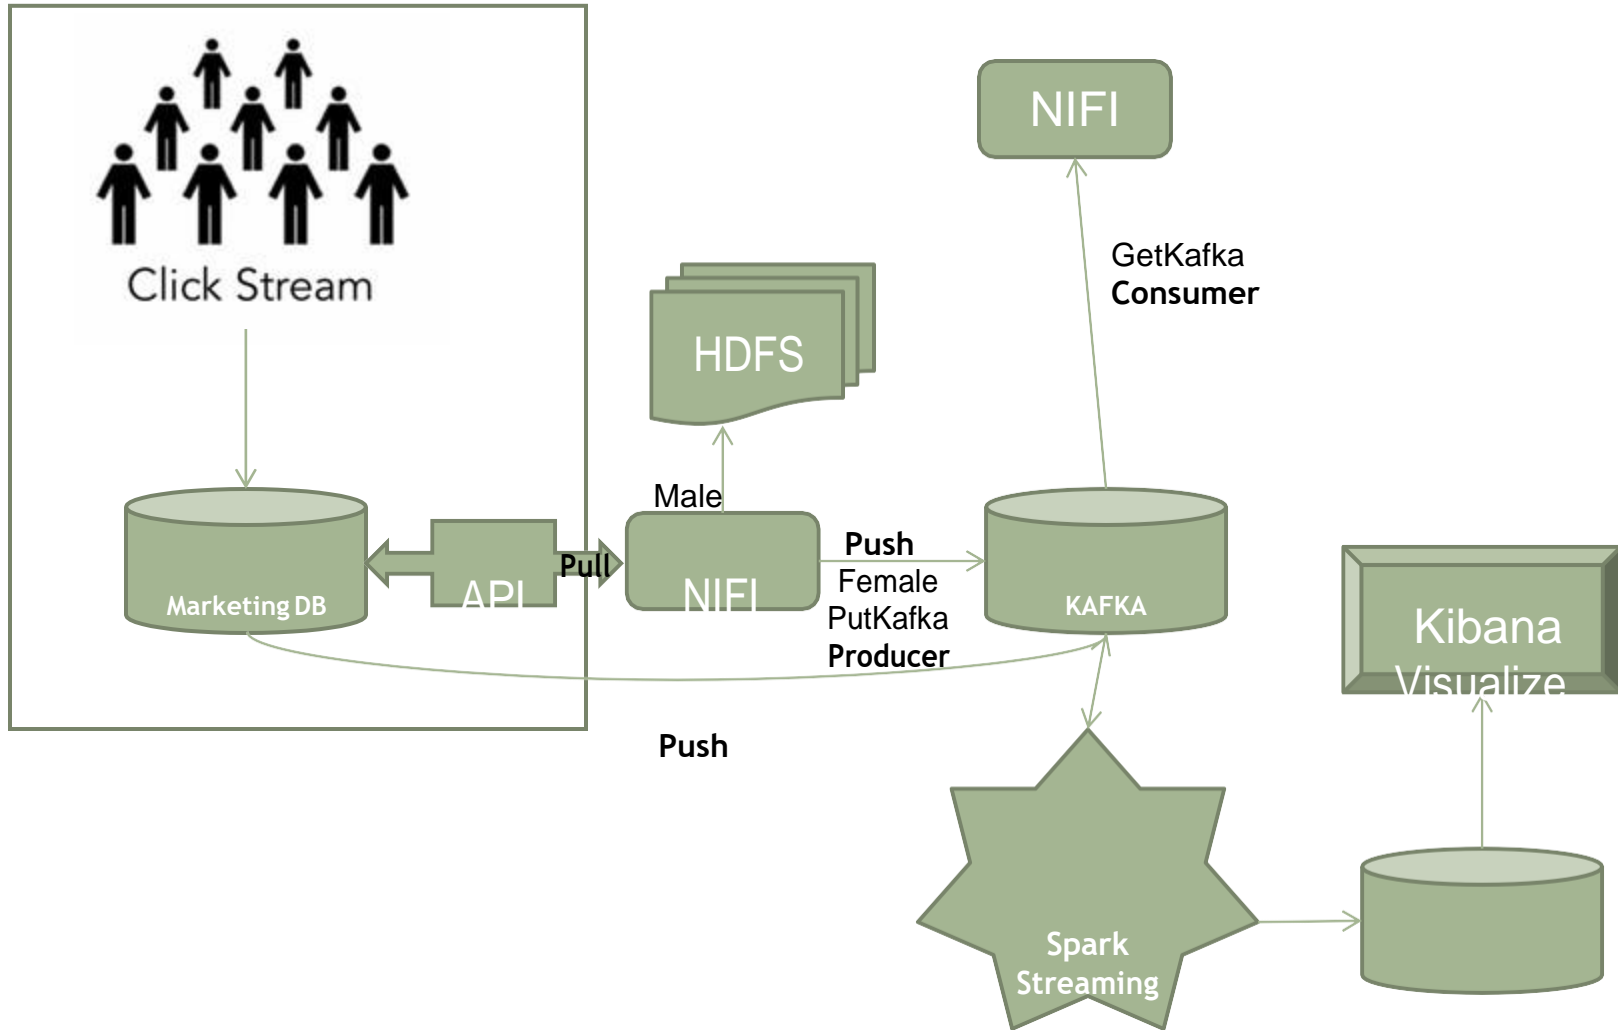


Click Stream

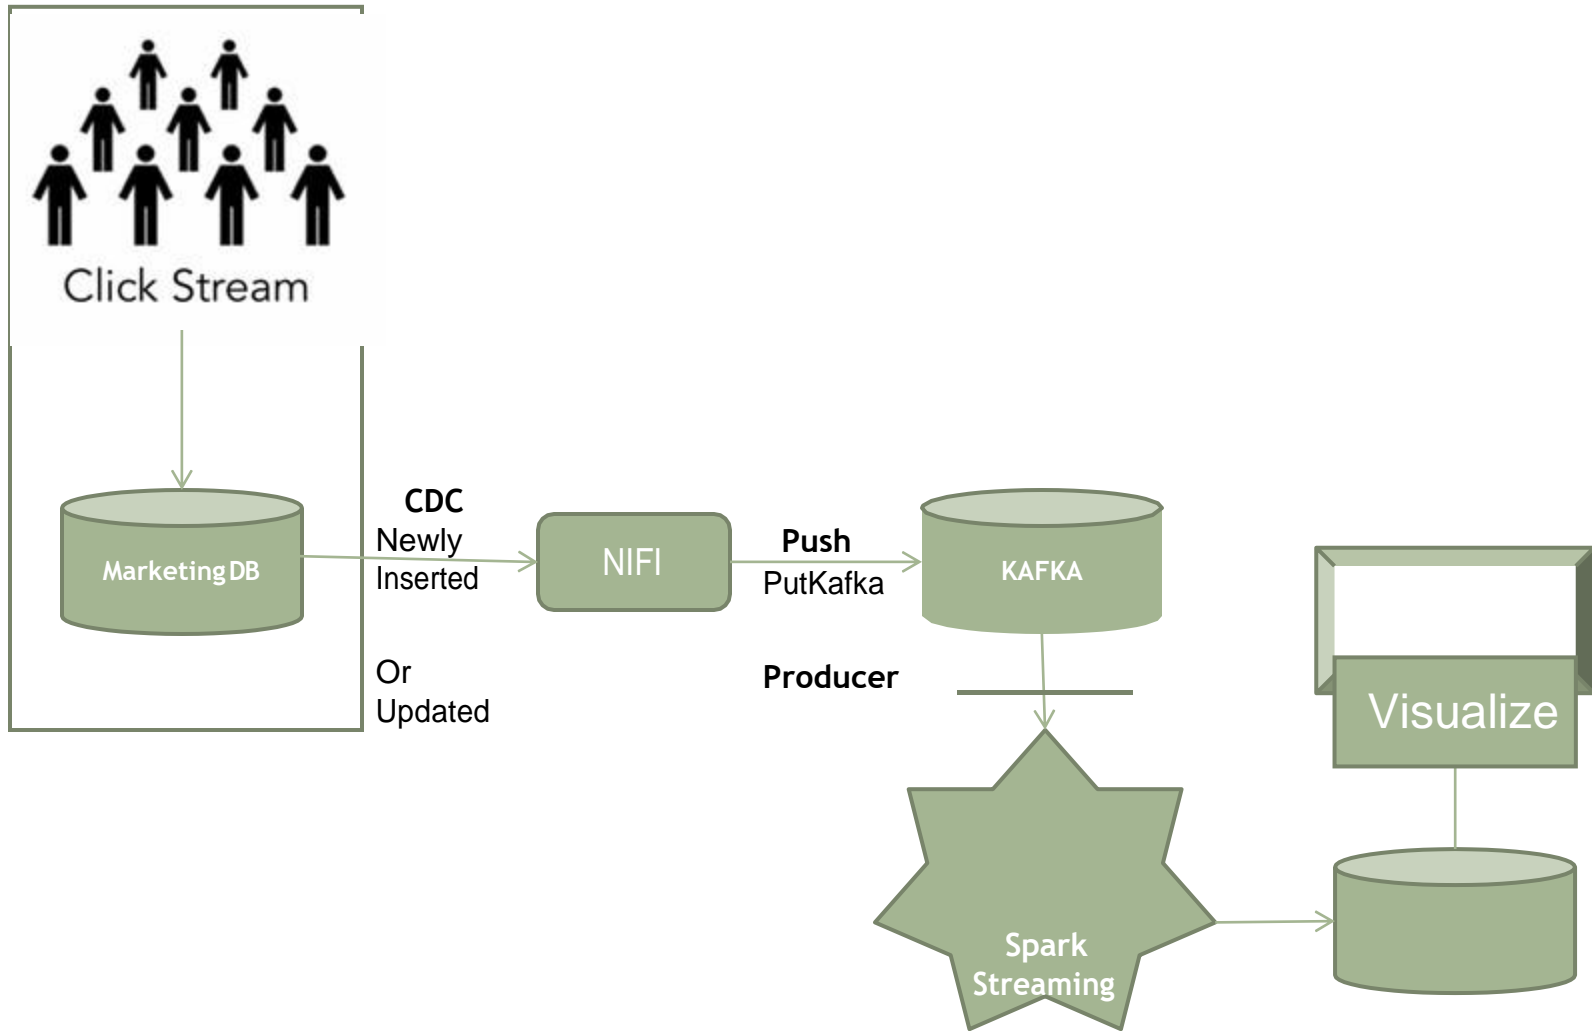
Fast Data Processing



Realtime Data Pipeline — Acquisition, Queue, Processing, Storage, Visualization

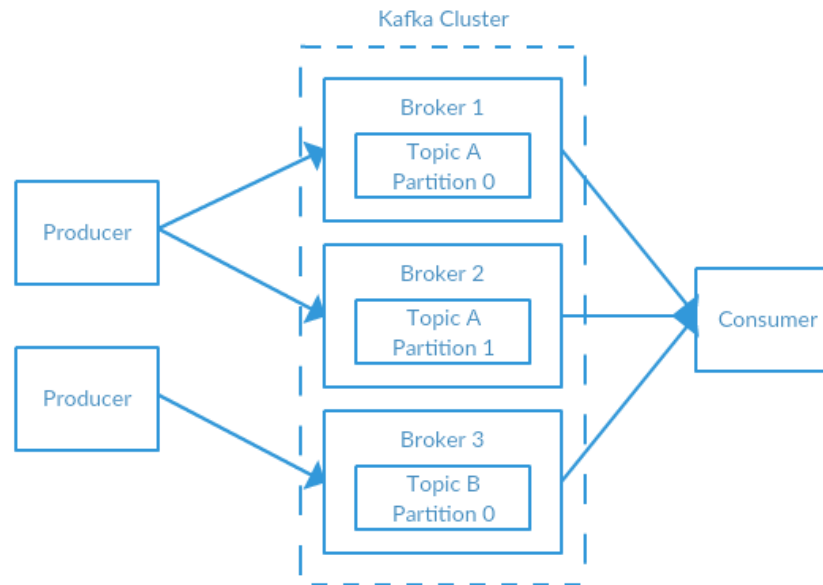


Realtime Data Pipeline



Components

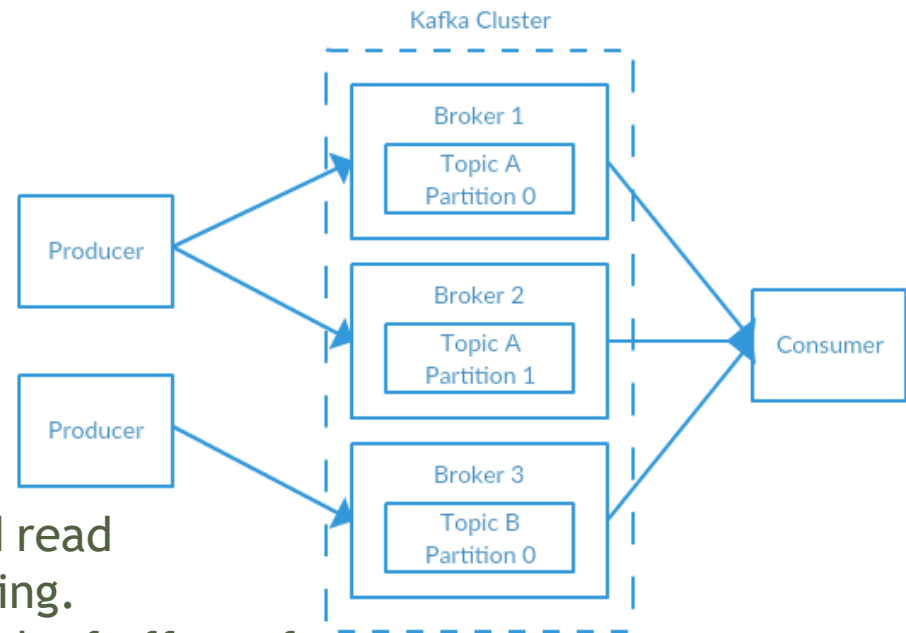
- Messages
- Producer
- Consumer
- Broker
- Kafka Cluster
- Topic
- Partition
- Replicas



HDFS	Kafka	SQL
File System	Message Queue	Database
Any Data	Any Data	Only Structured
Distribution & Scalable	Distributed & Scalable	No Distribution & Limited Scalability
File	Topic	Table
Blocks	Partition	Partition
Replica	Replica	No Replica
Fault Tolerance	Fault Tolerant	No Fault Tolerance
Batch	Message Queue	Database
Permenant long time data store	Temporary (staging)	Only Structured
No Incremental data management using offset	Offset for incremental data management	
Large sized file of small in numbers	Small messages of large in numbers	
	Used for Messaging System	
Used for Data Lake		

Components (Contd)

- **Messages** : The unit of data within Kafka with optional byte header, eg. Record in a table.
- **Broker - Managed by BigData Platform/Infra Team**
 - A single Kafka server
 - The broker receives messages from producers, assigns offsets to them, and commits on disk.
 - The broker serves consumer responding with the messages based on the partitions.
- **Producer**
 - Creates Messages.
 - Maintains messagekey to deliver messages to the particular partitions
- **Consumer**
 - Consume Messages.
 - Subscribe to one or more topic and read messages for processing or consuming.
 - Broker and Consumer can Keep track of offset of last consumed messages.



Components (Contd)

➤ Kafka Cluster

- Group of Brokers
- One Broker will act as a cluster Controller and assigns leader for the partition.

➤ Topic

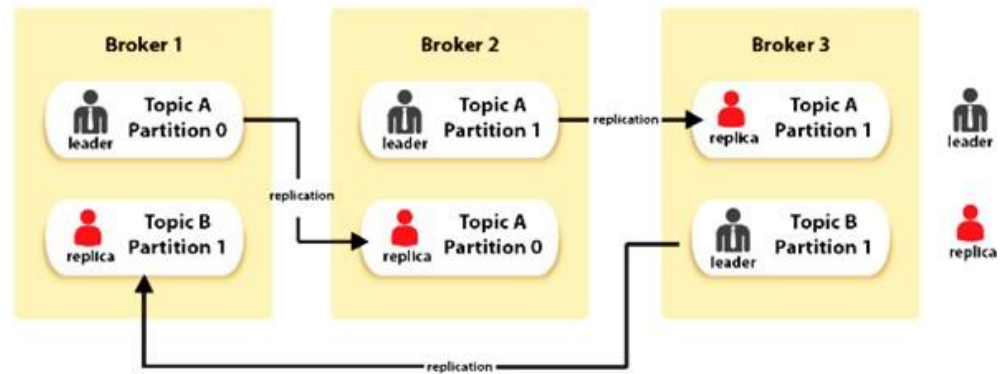
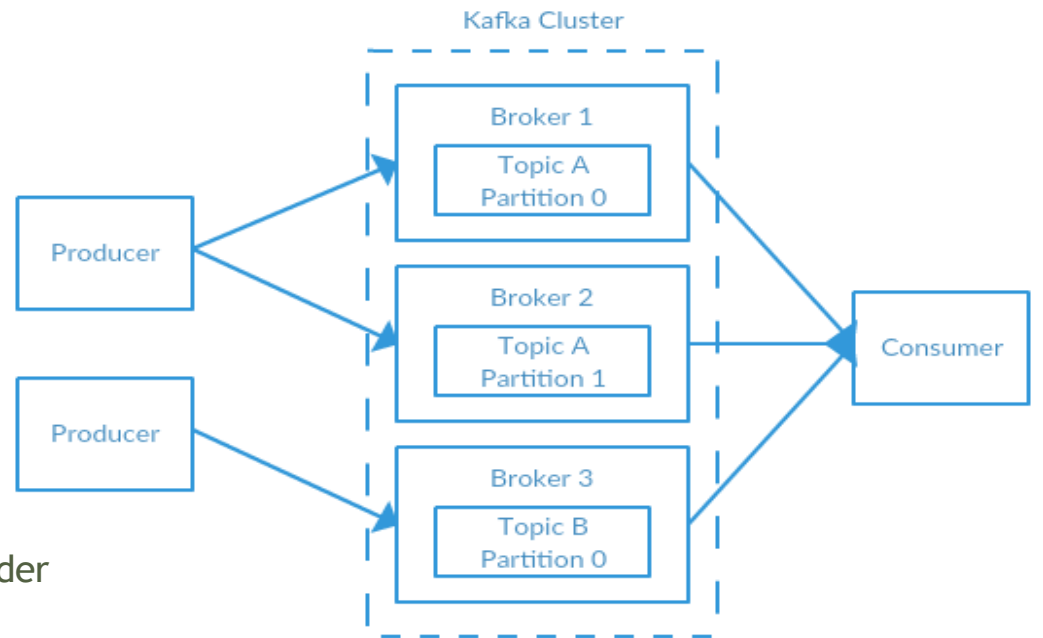
- Category of messages eg. Table/Folder
- Spread across all Brokers/nodes.

➤ Partition

- Horizontal division of topics.
- Scaled across multiple broker nodes.

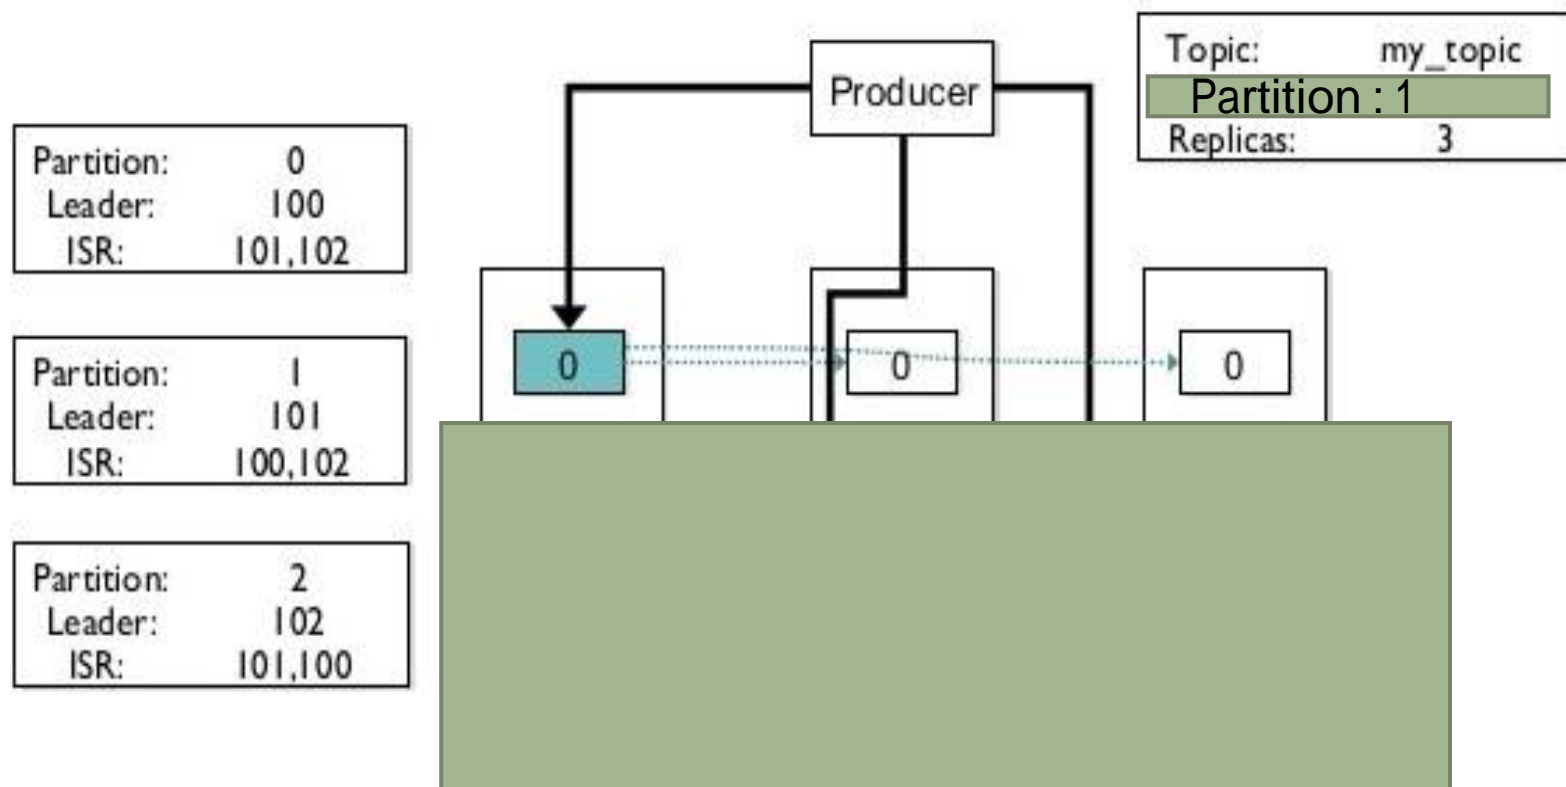
➤ Replicas

- Partitions in the Topics are replicated.
- Producer - ISR maintains - Fire and Forget, Sync and Async.



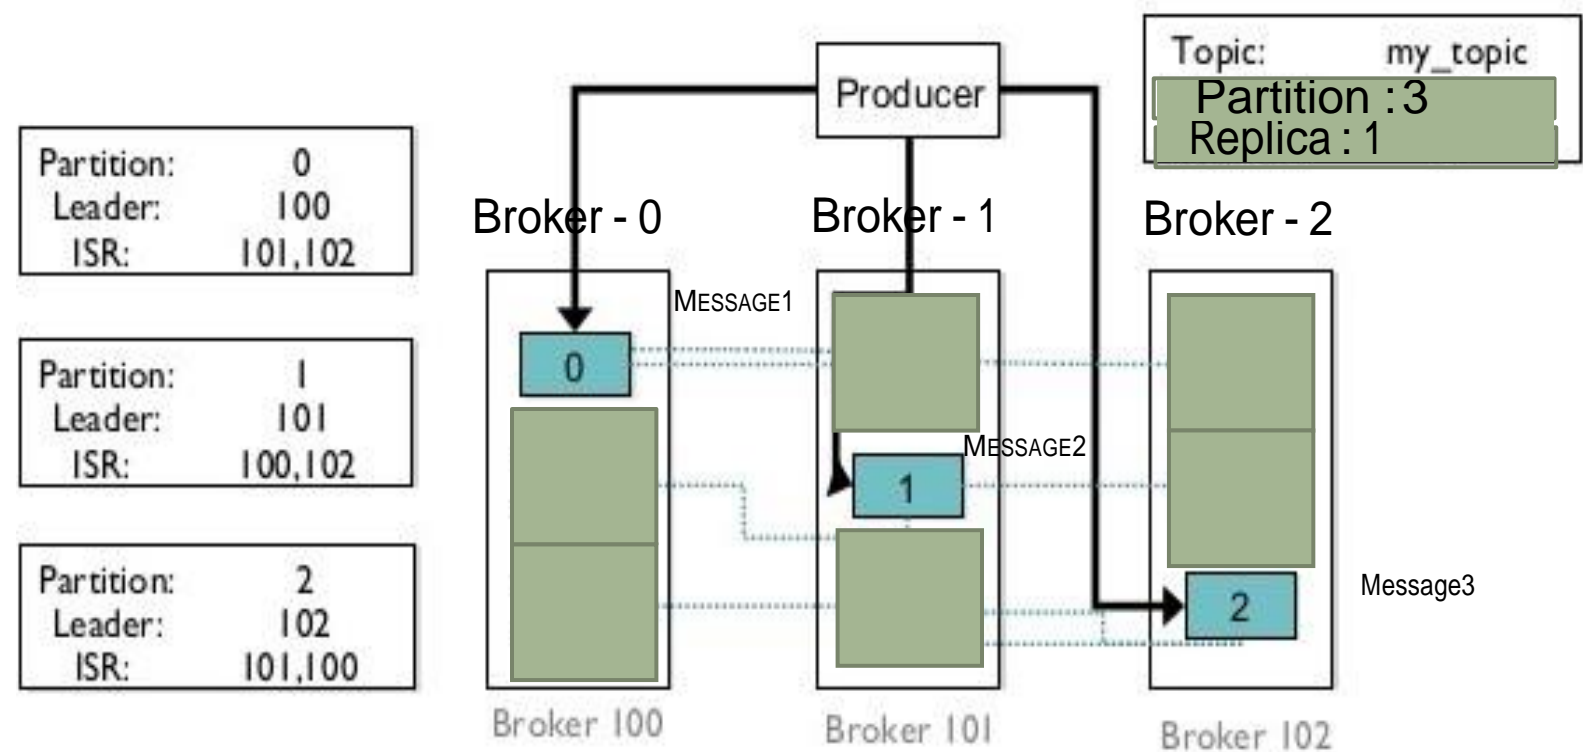
Replicas & ISR

Replication and ISR



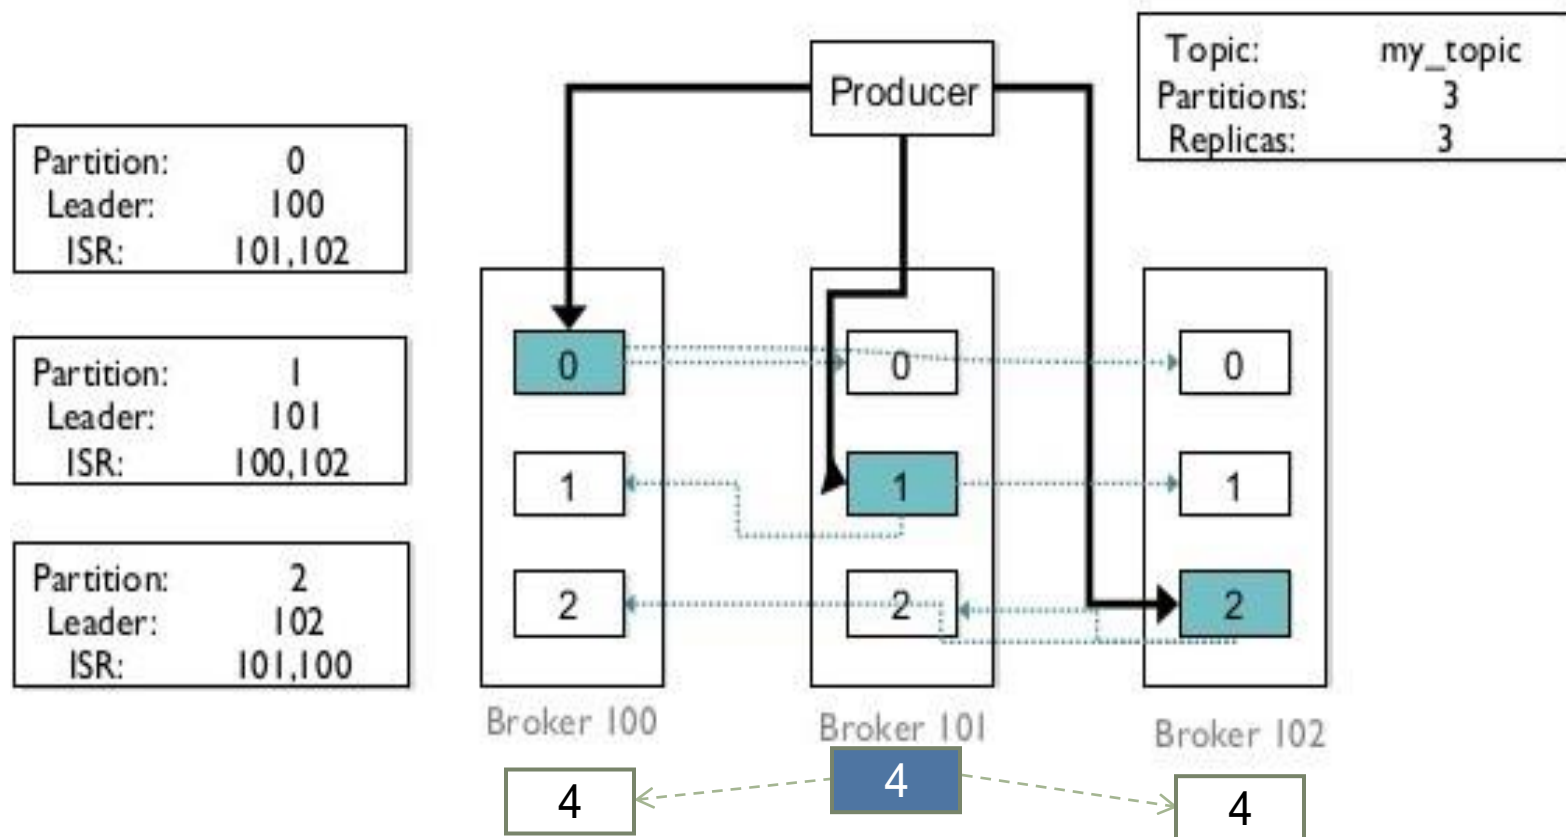
Replicas & ISR

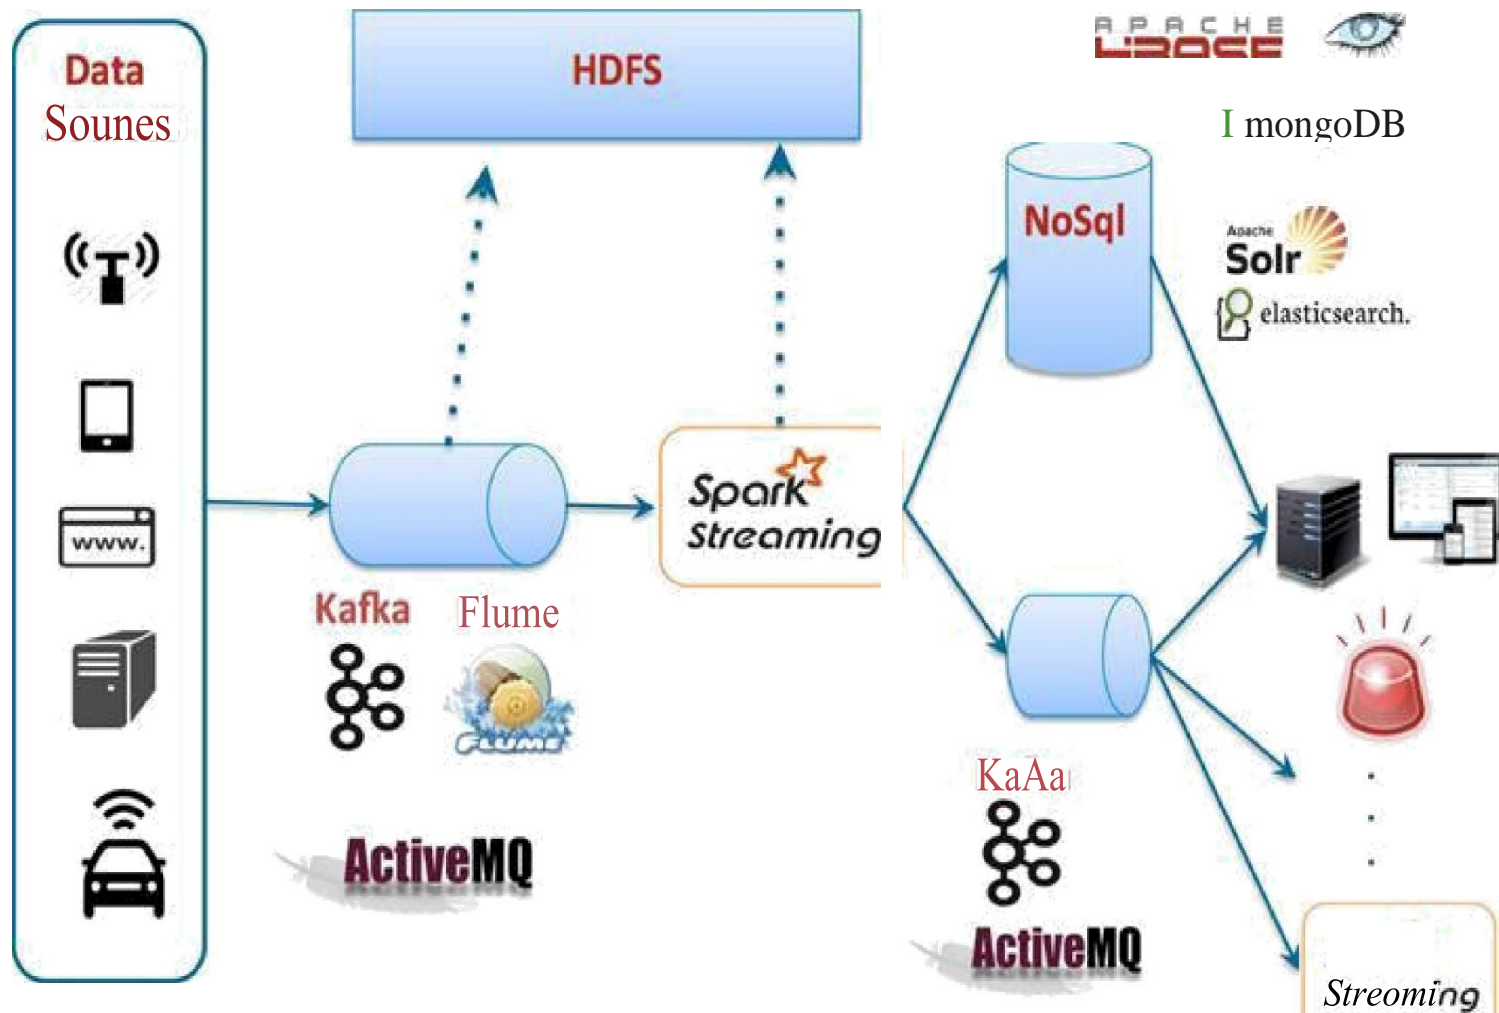
Replication and ISRs



Replicas & ISR

Replication and ISR





KAFKAWorkouts

1) Start the Zookeeper Coordination service:

```
zookeeper-server-start.sh -daemon /usr/local/kafka/config/zookeeper.properties
```

2) Start the Kafka server:

```
kafka-server-start.sh -daemon /usr/local/kafka/config/server.properties
```

Workouts:

Create a topic (with one replica and one partition)

```
kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --  
topic cts_topic
```

We can now see that topic if we run the list topic command:

```
kafka-topics.sh --list --bootstrap-server localhost:9092
```

Produce messages using Producer:

Kafka comes with a command line client that will take input from a file or from standard input and send it out as messages to the Kafka cluster. By default each line will be sent as a separate message. Run the producer and then type a few messages into the console to send to the server

```
kafka-console-producer.sh --broker-list localhost:9092 --topic cts_topic
```

Start a Consumer

Kafka also has a command line consumer that will dump out messages to standard output.

```
kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic cts_topic --from-beginning
```

Only look at the incremental logs

```
kafka-console-consumer.sh --zookeeper localhost:9092 --topic cts_topic
```

If you have each of the above commands running in a different terminal then you should now be able to type messages into the producer terminal and see them appear in the consumer terminal.

Setting up a multi-broker cluster

So far we have been running against a single broker. For Kafka, a single broker is just a cluster of size one, so nothing much changes other than starting a few more broker instances. But just to get feel for it, let's expand our cluster to three nodes (still all on our local machine).

First we make a config file for each of the brokers:

```
cp config/server.properties config/server-1.properties  
cp config/server.properties config/server-2.properties
```

Now edit these new files and set the following properties:

The broker.id property is the unique and permanent name of each node in the cluster. We have to override the port and log directory only because we are running these all on the same machine and we want to keep the brokers from all trying to register on the same port or overwrite each others data.

config/server-1.properties

=====

listeners=PLAINTEXT://:9093

broker.id=1 port=9093

log.dir=/tmp/kafka-logs-1

config/server-2.properties

=====

listeners=PLAINTEXT://:9094

broker.id=2 port=9094

log.dir=/tmp/kafka-logs-2

need to start the two new nodes:

kafka-server-start.sh -daemon \$KAFKA_HOME/config/server-1.properties

kafka-server-start.sh -daemon \$KAFKA_HOME/config/server-2.properties

Now create a new topic with a replication factor of three with a single partition:

```
kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 3 --partitions 1 --topic clickstream
```

```
kafka-topics.sh --list --bootstrap-server localhost:9092
```

Want to alter partitions use --alter

```
kafka-topics.sh --bootstrap-server localhost:9092 --alter --topic clickstream --partitions 4
```

How can we know which broker is doing what? To see that run the "describe topics" command:

```
kafka-topics.sh --describe --bootstrap-server localhost:9092 --topic clickstream
```

We can run the same command on the original topic we created to see where it is:

```
kafka-topics.sh --describe --bootstrap-server localhost:9092 --topic cts_topic
```

Let's publish a few messages to our new topic clickstream:

```
kafka-console-producer.sh --broker-list localhost:9092 --topic clickstream
```

Now let's consume these messages:

```
kafka-console-consumer.sh --bootstrap-server localhost:9092 --topic clickstream --from-beginning
```

Leadership has switched to one of the slaves and node 1 is no longer in the in-sync replica set:

```
kafka-topics.sh --describe --bootstrap-server localhost:9092 --topic clickstream
```

But the messages are still be available for consumption even though the leader that took the writes originally is down:

```
kafka-console-consumer.sh --bootstrap-server localhost:9092 --from-beginning --topic clickstream
```

If you want to delete the topic (mark for deletion, at the time of flush will be deleted else use delete.topic.enable=true)

```
kafka-topics.sh --delete --bootstrap-server localhost:9092 --topic cts_topic
```